

第十四讲：进入优化器和优化器追踪日志实现

知春路遇上八里桥

<2024-06-23 Sun>



- 1 前情提要
- 2 进入优化阶段
- 3 优化器追踪日志
- 4 MySQL 执行追溯日志




1

前情提要



执行流程



本节内容

• 连接器

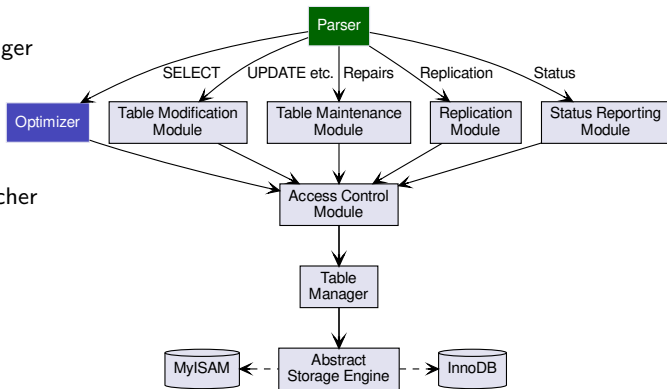
- ▶ ☒ 连接管理器 Connection Manager
- ▶ ☒ 线程管理器 Thread Manager
- ▶ ☒ 用户模块 User Module

• 解析器

- ▶ ☒ 网络模块 Net Module
- ▶ ☒ 派发模块 Commander Dispatcher
- ▶ ☒ 词法分析 Lexical Analysis
- ▶ ☒ 语法分析 Syntax Analysis

• 优化器

- ▶ ☒ 准备模块 Prepare Module
- ▶ ☐ 追踪日志 Optimizer Trace



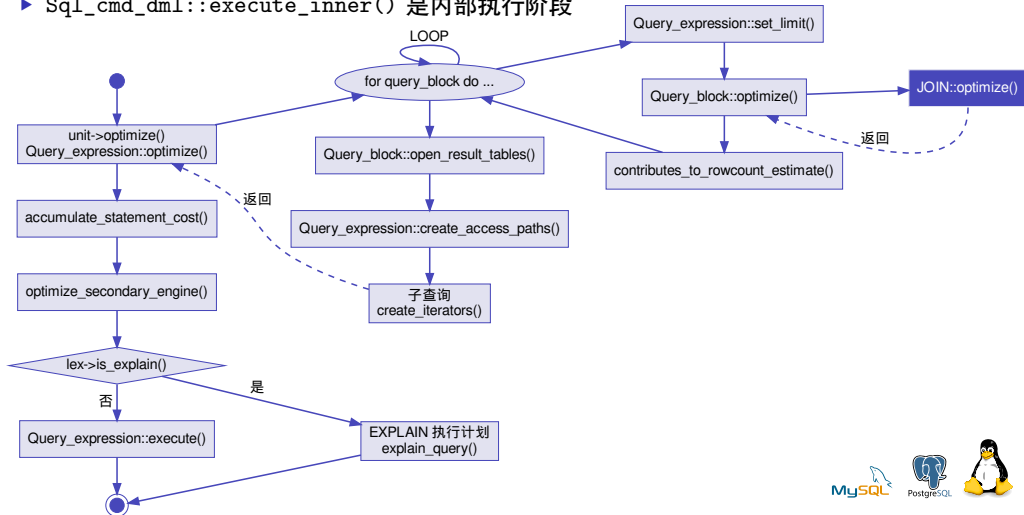
2

进入优化阶段



从 Sql_cmd_dml::execute_inner() 到 JOIN::optimize()

- Sql_cmd_dml::execute() 是 SELECT 执行的入口函数
 - Sql_cmd_dml::prepare() 是准备阶段，前一讲已经介绍
 - Sql_cmd_dml::execute_inner() 是内部执行阶段



优化器核心处理函数

- JOIN::optimize() 是优化阶段的入口函数
 - ▶ 它主要对 Query Block 进行优化
- JOIN::optimize() 函数将 Query_block 优化成 QEP

▶ `sql/sql_optimizer.cc`

```
337  bool JOIN::optimize(bool finalize_access_paths) {
338      DEBUG_TRACE;
339
340      uint no_jbuf_after = UINT_MAX;
341      Query_block *const set_operand_block =
342          query_expression()->non_simple_result_query_block();
343
344      assert(query_block->leaf_table_count == 0 ||
345             thd->lex->is_query_tables_locked() ||
346             query_block == set_operand_block);
347      assert(tables == 0 && primary_tables == 0 && tables_list == (Table_ref *)1);
348      :
1102     set_plan_state(ZERO_RESULT);
1103     return false;
1104 }
```

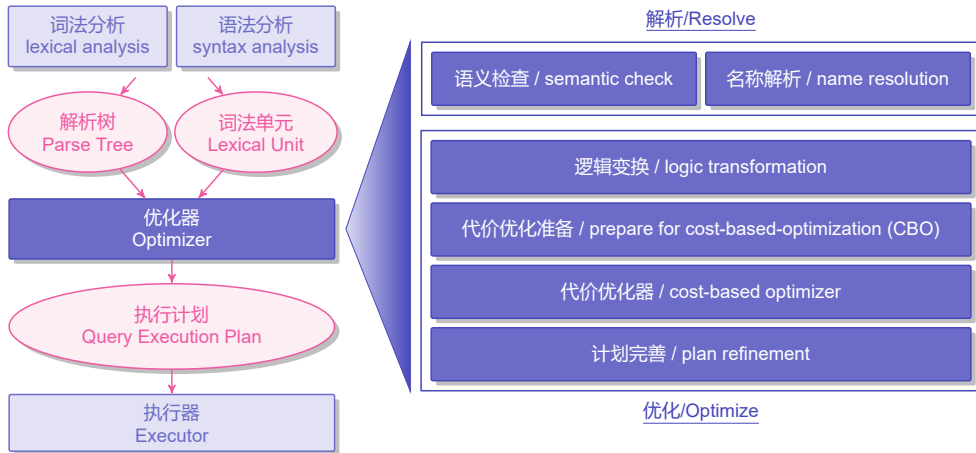


JOIN::optimize() 注释中的功能点

- # Logical transformations:
 - Outer to inner joins transformation.
 - Equality/constant propagation.
 - Partition pruning.
 - COUNT(*), MIN(), MAX() constant substitution in case of implicit grouping.
 - ORDER BY optimization.
- # Perform cost-based optimization of table order and access path selection.
See JOIN::make_join_plan()
- # Post-join order optimization:
 - Create optimal table conditions from the where clause and the join conditions.
 - Inject outer-join guarding conditions.
 - Adjust data access methods after determining table condition (several times.)
 - Optimize ORDER BY/DISTINCT.
- # Code generation
 - Set data access functions.
 - Try to optimize away sorting/distinct.
 - Setup temporary table usage for grouping and/or sorting.



优化器功能点

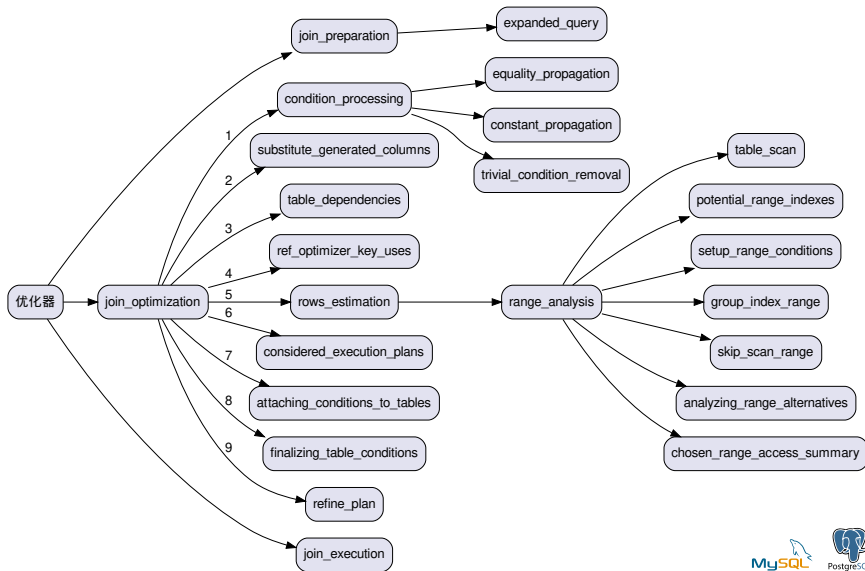


3

优化器追踪日志



Opt_trace 日志结构



源代码和 Opt_trace 日志 JSON 输出

- 通过代码搜索 "considered_execution_plans" 字符串

```
mysql-server $ grep -nR '"considered_execution' sql
sql/sql_planner.cc:2004:         "considered_execution_plans",
```

- 原始代码 `sql/sql_planner.cc`

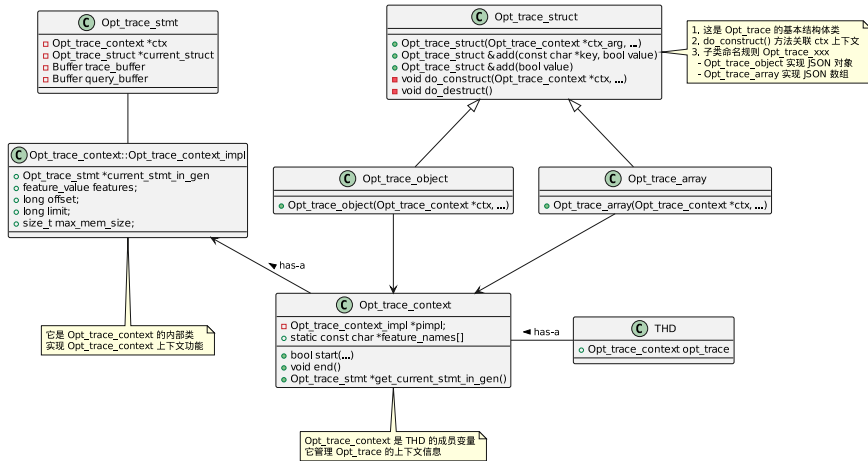
```
2002 Opt_trace_object wrapper(&join->thd->opt_trace);
2003 Opt_trace_array trace_plan(&join->thd->opt_trace,
2004                             "considered_execution_plans",
2005                             Opt_trace_context::GREEDY_SEARCH);
2006
2007 if (thd->optimizer_switch_flag(OPTIMIZER_SWITCH_COND_FANOUT_FILTER) &&
2008     join->where_cond) {
2009     for (uint idx = join->const_tables; idx < join->tables; ++idx)
2010         bitmap_clear_all(&join->best_ref[idx]->table()->cond_set);
2011
2012     /*
2013      * Set column bits for all columns involved in predicates in
2014      * cond_set. Used to avoid calculating condition filtering in
2015      * best_access_path() et al. when no filtering effect is possible.
2016      */
2017     join->where_cond->walk(&Item::add_field_to_cond_set_processor,
2018                           enum_walk::POSTFIX, nullptr);
2019 }
```

- 对应的追溯日志 `≈ OPT/Trace`

```
{
  "considered_execution_plans": [
    {
      "plan_prefix": [],
      "table": "`employees`",
      "best_access_path": {
        "considered_access_paths": [
          {
            "rows_to_scan": 9,
            "access_type": "range",
            "range_details": {
              "used_index": "PRIMARY"
            },
            "resulting_rows": 9,
            "cost": 2.81039,
            "chosen": true
          }
        ]
      },
      "condition_filtering_pct": 100,
      "rows_for_plan": 9,
      "cost_for_plan": 2.81039,
      "chosen": true
    }
  ]
}
```



Opt_trace 追踪优化器



(gdb) set print elements 400

(gdb) p thd->opt_trace.get_current_stmt_in_gen()->trace_buffer.c_ptr_safe()

```
$83 = 0x7fff3019dde0 {"\n  \"steps\": [\n    {\n      \"join_preparation\": {\n        \"select#\": 1,\n        \"steps\": [\n          \"expanded_query\": \"/* select#1 */ select `employees`.`emp_no` AS `emp_no`,`em...
```



4

MySQL 执行追溯日志



原始日志分析 - 壹

- T08 其中 8 是连接 ID (Connection id)
- 每个函数的调用开始和结束通过 < 和 > 包裹
- 函数调用栈通过 | 缩进显示
- 中间夹杂着 DEBUG_PRINT("info", ...) 的输出

```
T08: | | | >lex_start
T08: | | | >alloc_query
T08: | | | | thd_query: thd->thread_id():8 thd:0x7fff30001050 query:SELECT * FROM employees WHERE emp_no > 1234 LIMIT 1
T08: | | | >parse_sql
T08: | | | >mysql_execute_command
T08: | | | | >Table_ref*, enum_sql_command, List<set_var_base>*, const char*, size_t, sp_printable*, const CHARSET_INFO*
T08: | | | | | >Opt_trace_context::start
T08: | | | | | | opt: new stmt 0x7fff3000fe20 support_I_S 0
T08: | | | | | | >Opt_trace_context::purge_stmts
T08: | | | | | | <Opt_trace_context::purge_stmts
T08: | | | | | | opt: rc 1048576 max_mem_size 1048576
T08: | | | | | | <Opt_trace_context::start
T08: | | | | | | >{anonymous}::opt_trace_disable_if_no_tables_access
T08: | | | | | | <{anonymous}::opt_trace_disable_if_no_tables_access
T08: | | | | | >open_temporary_tables
T08: | | | | | >open_temporary_table
T08: | | | | | | enter: table: 'employees'.'employees'
T08: | | | | | <open_temporary_table
T08: | | | | | <open_temporary_tables
```



原始日志分析 - 贰

● 进入优化器之前的准备阶段 Query_block::prepare()

```
T08: | | | | >bool Sql_cmd_dml::execute
T08: | | | | | >bool Sql_cmd_dml::prepare
T08: | | | | | | >check_table_access
T08: | | | | | | | info: table: employees derived: 0 view: 0
T08: | | | | | | | >check_access
T08: | | | | | | | | enter: db: employees want_access: 1 master_access: 2147483647
T08: | | | | | | | | THD::enter_stage: 'checking permissions' /opt/src/mysql-server/sql/auth/sql_authorization.cc:2146
T08: | | | | | | | >open_tables_for_query
T08: | | | | | | | >open_tables
T08: | | | | | | | | THD::enter_stage: 'Opening tables' /opt/src/mysql-server/sql/sql_base.cc:5797
T08: | | | | | | | <open_tables
T08: | | | | | | | >Query_block::prepare
T08: | | | | | | | | opt: (null): starting struct
T08: | | | | | | | | opt: join_preparation: starting struct
T08: | | | | | | | | opt: select#: 1
T08: | | | | | | | | opt: steps: starting struct
T08: | | | | | | | >Query_block::setup_tables
T08: | | | | | | | <Query_block::setup_tables
T08: | | | | | | | >Query_block::setup_wild
T08: | | | | | | | <Query_block::prepare
T08: | | | | | <bool Sql_cmd_dml::prepare
T08: | | | | | THD::enter_stage: 'init' /opt/src/mysql-server/sql/sql_select.cc:772
```



原始日志分析 - 叁

• optimize_cond() 函数执行的细节

```
T08: | | | | | >Query_expression::optimize
T08: | | | | | | >Query_block::optimize
T08: | | | | | | | >JOIN::optimize
T08: | | | | | | | THD::enter_stage: 'optimizing' /opt/src/mysql-server/sql/sql_optimizer.cc:354
T08: | | | | | | | opt: (null): starting struct
T08: | | | | | | | opt: join_optimization: starting struct
T08: | | | | | | | opt: select#: 1
T08: | | | | | | | opt: steps: starting struct
T08: | | | | | | | >optimize_cond
T08: | | | | | | | | opt: (null): starting struct
T08: | | | | | | | | opt: condition_processing: starting struct
T08: | | | | | | | | opt: condition: "WHERE"
T08: | | | | | | | | opt: original_condition: "(`employees`.`emp_no` > 1234)"
T08: | | | | | | | | opt: steps: starting struct
T08: | | | | | | | | opt: (null): starting struct
T08: | | | | | | | | opt: transformation: "equality_propagation"
T08: | | | | | | | | opt: subselect_evaluation: starting struct
T08: | | | | | | | | opt: subselect_evaluation: ending struct
T08: | | | | | | | | opt: resulting_condition: "(`employees`.`emp_no` > 1234)"
T08: | | | | | | | | opt: (null): ending struct
T08: | | | | | | | | opt: (null): starting struct
T08: | | | | | | | | opt: transformation: "constant_propagation"
T08: | | | | | | | | opt: subselect_evaluation: starting struct
T08: | | | | | | | | opt: subselect_evaluation: ending struct
T08: | | | | | | | | opt: resulting_condition: "(`employees`.`emp_no` > 1234)"
```



结束

