

# 第十六讲：联表查询 Join 优化器的设计与实现

知春路遇上八里桥

<2024-06-29 Sat>



## 1 前情提要

## 2 Join

## 3 代码分析




1

## 前情提要



# 执行流程



# 本节内容

## • 连接器

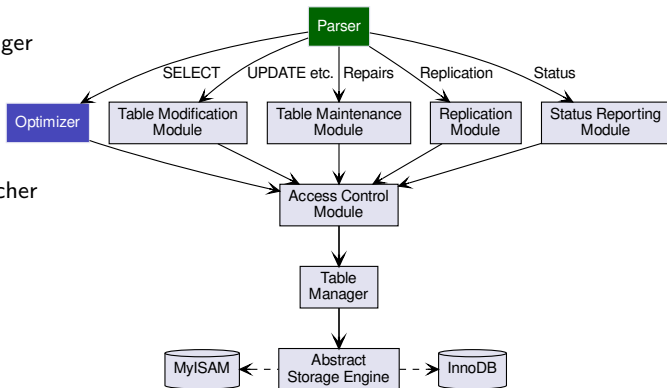
- ▶ ☒ 连接管理器 Connection Manager
- ▶ ☒ 线程管理器 Thread Manager
- ▶ ☒ 用户模块 User Module

## • 解析器

- ▶ ☒ 网络模块 Net Module
- ▶ ☒ 派发模块 Commander Dispatcher
- ▶ ☒ 词法分析 Lexical Analysis
- ▶ ☒ 语法分析 Syntax Analysis

## • 优化器

- ▶ ☒ 准备模块 Prepare Module
- ▶ ☒ 追踪日志 Optimizer Trace
- ▶ ☐ 优化模块 Optimize Module



2

Join



# Join 介绍

- 为了方便编写查询语句, MySQL 提供若干类型的 Join 语句<sup>1</sup>
- inner join 返回左表和右表都包含的行, 这里左表和右表的次序可以交换
- left join 和 right join 表的顺序不可以交换, 其中
  - ▶ left join 返回数据中左表必选, 右表可选
  - ▶ right join 返回数据中右表必选, 左表可选
- straight join 的逻辑和 join 一样, 只不过左表总是在右边之前读取
  - ▶ straight join 提供了控制执行器读取表先后顺序
  - ▶ 它的实现参考 Optimize\_table\_order::optimize\_straight\_join()
- 一些 Join 的例子

```
select * from t1, t2;  
select * from t1 inner join t2 on t1.id = t2.id;  
select * from t1 left join t2 on t1.id = t2.id left join t3 on t2.id = t3.id;  
select * from t1 left join t2 on t1.id = t2.id;  
select * from t1 left join t2 using (id);
```

<sup>1</sup><https://dev.mysql.com/doc/refman/8.0/en/join.html>

# 查询结果实例 (壹)

表 t1 和 t2 中的数据

```
mysql> select * from t1;
```

id	c
1	aaa
2	bbb

```
mysql> select * from t2;
```

id	c
1	xxx
3	yyy

- join 连接

```
mysql> select * from t1 join t2;
```

id	c	id	c
2	bbb	1	xxx
1	aaa	1	xxx
2	bbb	3	yyy
1	aaa	3	yyy

- 笛卡尔积

```
mysql> select * from t1, t2;
```

id	c	id	c
2	bbb	1	xxx
1	aaa	1	xxx
2	bbb	3	yyy
1	aaa	3	yyy

- straight join 直接连接

```
mysql> select * from t1 straight_join t2;
```

id	c	id	c
2	bbb	1	xxx
1	aaa	1	xxx
2	bbb	3	yyy
1	aaa	3	yyy

- inner join 内连接

```
mysql> select * from t1 inner join t2;
```

id	c	id	c
2	bbb	1	xxx
1	aaa	1	xxx
2	bbb	3	yyy
1	aaa	3	yyy



## 查询结果实例 (貳)

表 t1 和 t2 中的数据

```
mysql> select * from t1;
```

id	c
1	aaa
2	bbb

```
mysql> select * from t2;
```

id	c
1	xxx
3	yyy

- 内连接

```
mysql> select * from t1 inner join t2 on t1.id = t2.id;
```

id	c	id	c
1	aaa	1	xxx

- 左连接

```
mysql> select * from t1 left join t2 on t1.id = t2.id;
```

id	c	id	c
1	aaa	1	xxx
2	bbb	NULL	NULL


- 右连接

```
mysql> select * from t1 right join t2 on t1.id = t2.id;
```


id	c	id	c
1	aaa	1	xxx
NULL	NULL	3	yyy

# 连接树


- 多个表进行连接时，执行过程可以表示成树形结构，即连接树 (Join Tree)
  - 左深树 (Left Deep Join Tree) 的每个连接的右节点都是一个表
  - 右深树 (Right Deep Join Tree) 的每个连接的左节点都是一个表
  - 浓密树 (Bushy Join Tree) 的左节点或右节点都有可能不是表



左深树




右深树



浓密树

# 多表 Join 的左深树结构

- $n$  个表会有排列数  $^2 A_n^n$  种做 Join 情况. 对 a, b, c 三个表,  $A_3^3 = 6$  种情况
  - ▶ (a, b, c) / (a, c, b) / (b, a, c)
  - ▶ (b, c, a) / (c, a, b) / (c, b, a)
- MySQL 多表 Join 的查询语句采用数组来表示左深树结构



```
1 select
2   e.first_name, e.last_name, a.from_date
3 from
4   dept_manager a
5   join departments d on a.dept_no = d.dept_no
6   join employees e on a.emp_no = e.emp_no
7 where
8   d.dept_no = 'd001'
```

- 最终优化后的 Join 顺序为 d/a/e , 并非输入顺序 a/d/e

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	d	NULL	const	PRIMARY	PRIMARY	16	const	1	100.00	Using index
1	SIMPLE	a	NULL	ref	PRIMARY,dept_no	dept_no	16	const	2	100.00	Using index condition
1	SIMPLE	e	NULL	eq_ref	PRIMARY	PRIMARY	4	employees.a.emp_no	1	100.00	NULL

<sup>2</sup>排列数  $A_m^n$  表示从  $n$  个不同的元素中任取  $m$  ( $m \leq n$ ) 个所有排列的个数.  $A_m^n = n!/(n-m)!$



# Json 格式详细执行计划

```
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "2.90"
    },
    "nested_loop": [
      {
        "table": {
          "table_name": "d",
          "access_type": "const",
          "possible_keys": [
            "PRIMARY"
          ],
          "key": "PRIMARY",
          "used_key_parts": [
            "dept_no"
          ],
          "key_length": "16",
          "ref": [
            "const"
          ],
          "rows_examined_per_scan": 1,
          "rows_produced_per_join": 1,
          "filtered": "100.00",
          "using_index": true,
          "cost_info": {
            "read_cost": "0.00",
            "eval_cost": "0.10",
            "prefix_cost": "0.00",
            "data_read_per_join": "184"
          },
          "used_columns": [
            "dept_no"
          ]
        }
      },

```

```
    {
      "table": {
        "table_name": "a",
        "access_type": "ref",
        "possible_keys": [
          "PRIMARY",
          "dept_no"
        ],
        "key": "dept_no",
        "used_key_parts": [
          "dept_no"
        ],
        "key_length": "16",
        "ref": [
          "const"
        ],
        "rows_examined_per_scan": 2,
        "rows_produced_per_join": 2,
        "filtered": "100.00",
        "index_condition": "(`employees`.`a`.`dept_no` = 'd001')",
        "cost_info": {
          "read_cost": "0.50",
          "eval_cost": "0.20",
          "prefix_cost": "0.70",
          "data_read_per_join": "64"
        },
        "used_columns": [
          "emp_no",
          "dept_no",
          "from_date"
        ]
      },
    ]
  },

```

```
    {
      "table": {
        "table_name": "e",
        "access_type": "eq_ref",
        "possible_keys": [
          "PRIMARY"
        ],
        "key": "PRIMARY",
        "used_key_parts": [
          "emp_no"
        ],
        "key_length": "4",
        "ref": [
          "employees.a.emp_no"
        ],
        "rows_examined_per_scan": 1,
        "rows_produced_per_join": 2,
        "filtered": "100.00",
        "cost_info": {
          "read_cost": "2.00",
          "eval_cost": "0.20",
          "prefix_cost": "2.90",
          "data_read_per_join": "272"
        },
        "used_columns": [
          "emp_no",
          "first_name",
          "last_name"
        ]
      }
    ]
  },

```



## 另外两种 Join 的写法

- join3\_02 第二种 Join 写法: 将条件全部写在 where 的后面

```
1  select
2      e.first_name, e.last_name, a.from_date
3  from
4      dept_manager a join departments d join employees e
5  where
6      a.dept_no = d.dept_no and a.emp_no = e.emp_no and d.dept_no = 'd001'
```

- join3\_03 第三种 Join 写法: 笛卡尔积

```
1  select
2      e.first_name, e.last_name, a.from_date
3  from
4      dept_manager a, departments d, employees e
5  where
6      a.dept_no = d.dept_no and a.emp_no = e.emp_no and d.dept_no = 'd001'
```

- 后面两种写法和之前的会产生同样的执行计划<sup>3</sup>

<sup>3</sup><https://github.com/Jeanhwea/mysql-source-course/blob/master/assets/join3-prepare-example.org>

# 尝试一下 straight\_join 来控制 Join 顺序

- 使用 straight\_join 查询语句实验<sup>a</sup>

```
1 select
2     e.first_name,
3     e.last_name,
4     a.from_date
5 from
6     dept_manager a
7     straight_join departments d
8         on a.dept_no = d.dept_no
9     straight_join employees e
10        on a.emp_no = e.emp_no
11 where
12     d.dept_no = 'd001'
```

- 通过估算得到 cost 值是:  $3.35 > 2.90$

```
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "3.35"
    },
    "nested_loop": [
      {
        "table": {
          "table_name": "a",
          "access_type": "ref",
          "possible_keys": [
            "PRIMARY",
            "dept_no"
          ],
          "key": "dept_no",
          "used_key_parts": [
            "dept_no"
          ],
          "key_length": "16",
          "ref": [
            "const"
          ],

```

<sup>a</sup>见 JOIN\_ORDER Hint 的用法

# 对比 straight\_join 实际执行耗时

```
mysql> set profiling=1;  
Query OK, 0 rows affected, 1 warning (0.00 sec)
```

```
mysql> ... 执行查询
```

```
mysql> show profiles\G  
***** 1. row *****  
Query_ID: 1  
Duration: 0.00922100  
  Query: select /*+ rtc */  
    e.first_name, e.last_name, a.from_date  
  from  
    dept_manager a straight_join departments d straight_join employees e  
  where  
    a.dept_no = d.dept_no and a.emp_no = e.emp_no and d.dept_no = 'd001'  
***** 2. row *****  
Query_ID: 2  
Duration: 0.00867425  
  Query: select /*+ rtc */  
    e.first_name, e.last_name, a.from_date  
  from  
    dept_manager a join departments d join employees e  
  where  
    a.dept_no = d.dept_no and a.emp_no = e.emp_no and d.dept_no = 'd001'  
2 rows in set, 1 warning (0.00 sec)
```



# Join 执行算法

- MySQL 执行器默认使用 NLJ<sup>4</sup> 算法来实现 Join

- ▶ 假设有三个表 t1, t2 和 t3 对于的 Join 类型如下

- ① t1 range 范围扫描
- ② t2 ref 引用
- ③ t3 ALL 全表扫描

- ▶ NLJ (Nested-Loop Join) 算法的伪代码如下

```
for each row in t1 matching range {  
    for each row in t2 matching reference key {  
        for each row in t3 {  
            if row satisfies join conditions, send to client  
        }  
    }  
}
```

- NLJ 的思路是：先选定驱动表，然后根据过滤条件来查询被驱动表的数据

- ▶ 如果过滤条件在 Join 操作之前，NLJ 算法的性能很好
- ▶ 通常经过优化器后会满足小表驱动大表的原则

<sup>4</sup><https://dev.mysql.com/doc/refman/8.0/en/nested-loop-joins.html>



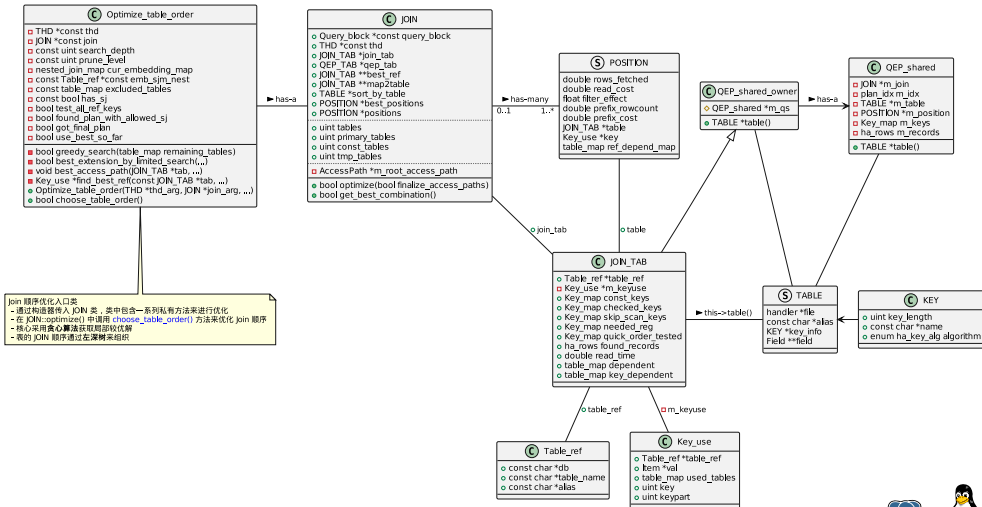


3

## 代码分析



# Join Order 数据结构



# 条件断点调试过程

## ❶ 设置条件断点，如果包含 "rtc" 字符串则触发断点

```
(gdb) break JOIN::optimize if strstr(thd->query().str, "rtc")  
(gdb) c
```

## ❷ 添加标识 `/*+ rtc */` 后发送给 mysqld

```
(gdb) printf "%s\n", thd->query().str  
select /*+ rtc */  
    e.first_name, e.last_name, a.from_date  
from  
    dept_manager a join departments d join employees e  
where  
    a.dept_no = d.dept_no and a.emp_no = e.emp_no and d.dept_no = 'd001'  
(gdb)
```



## Optimize\_table\_order::choose\_table\_order()

- choose\_table\_order() 选择多表 Join 的顺序
- greedy\_search() 穷举式 + 贪婪式算法, 在尽可能小的搜索空间内得到相对较优的查询计划
  - ▶ 算法输入 remaining\_tables 剩余待计算的表
  - ▶ 算法输出 pplan (partial plan) 当前最优的计划
  - ▶ 算法实现的伪代码如下

```
1  pplan = <>;
2  do {
3      (t, a) = best_extension(pplan, remaining_tables);
4      pplan = concat(pplan, (t, a));
5      remaining_tables = remaining_tables - t;
6  } while (remaining_tables != {});
7  return pplan;
```

- best\_extension\_by\_limited\_search() 最终找到近似最优解的连接排列组合
  - ▶ 对于  $n$  个表的搜索量为  $A_n^n$  整体搜索空间比较大, 所以会进行剪枝
    - ① pruned\_by\_cost: 根据代价剪枝
    - ② pruned\_by\_heuristic: 根据启发式规则剪枝
  - ▶ 限制最大可搜索的深度<sup>5</sup>, 记录在 Optimize\_table\_order::search\_depth 中
- best\_access\_path() 获取最优路径

<sup>5</sup>最大搜索深度参数 optimizer\_search\_depth=62



# 贪心算法递归搜索过程

## • 包含 5 个表的 Join 搜索过程

```
1 select /*+ rtc */
2     d.dept_no, t.title, e.first_name, s.from_date, s.salary
3 from
4     dept_manager a join departments d join employees e join salaries s join titles t
5 where
6     a.dept_no = d.dept_no and a.emp_no = e.emp_no and a.emp_no = s.emp_no
7     and a.emp_no = t.emp_no and d.dept_no = 'd001';
```

## • 递归调用过程，当 pplan\_cost 值过大时发生剪枝<sup>6</sup>

```
$ grep cost_for_plan /tmp/mysqld.trace
T@10: | | opt: cost_for_plan: 0.461096          => POSITIONS: d          <= 搜索
T@10: | | | opt: cost_for_plan: 0.451096        => POSITIONS: d a        <= 搜索
T@10: | | | | opt: cost_for_plan: 2.64926       => POSITIONS: d a e      <= 搜索
T@10: | | | | | opt: cost_for_plan: 4.99109     => POSITIONS: d a e t    <= 搜索
T@10: | | | | | | opt: cost_for_plan: 10.7898   => POSITIONS: d a e t s  <= 目前为止最优解
T@10: | | | | | | | opt: cost_for_plan: 6.58652 => POSITIONS: d a e s    <= 剪枝
T@10: | | | | | | | opt: cost_for_plan: 2.79293 => POSITIONS: d a t      <= 剪枝
T@10: | | | | | | | opt: cost_for_plan: 4.38835 => POSITIONS: d a s      <= 剪枝
T@10: | | | | | | | opt: cost_for_plan: 28466.5 => POSITIONS: d e        <= 剪枝
T@10: | | | | | | | opt: cost_for_plan: 44874.5 => POSITIONS: d t        <= 剪枝
T@10: | | | | | | | opt: cost_for_plan: 281270  => POSITIONS: d s        <= 剪枝
```

<sup>6</sup><https://github.com/Jeanhwea/mysql-source-course/blob/master/assets/join5.mysqld.trace>



## Optimize\_table\_order::find\_best\_ref()

- find\_best\_ref() 获取最优的索引来进行 ref 访问，使用以下优先级进行选取
  - 1) A clustered primary key with equality predicates on all keyparts is always chosen.
  - 2) A non nullable unique index with equality predicates on all keyparts is preferred over a non-unique index, nullable unique index or unique index where there are some keyparts without equality predicates.
  - 3) Otherwise, the index with best cost estimate is chosen.
- 索引选取优先级
  - ▶ 等值聚簇索引
  - ▶ 非空唯一等值索引
  - ▶ 非空普通等值索引
  - ▶ 可空唯一等值索引
  - ▶ 唯一非等值索引
  - ▶ Cost 最小的索引



# 结束

