

# 第十七讲：子查询优化器的设计与实现

知春路遇上八里桥

<2024-07-04 Thu>



## 1 前情提要

## 2 子查询

## 3 半连接

## 4 其他类型




1

## 前情提要



# 执行流程



# 本节内容

## • 连接器

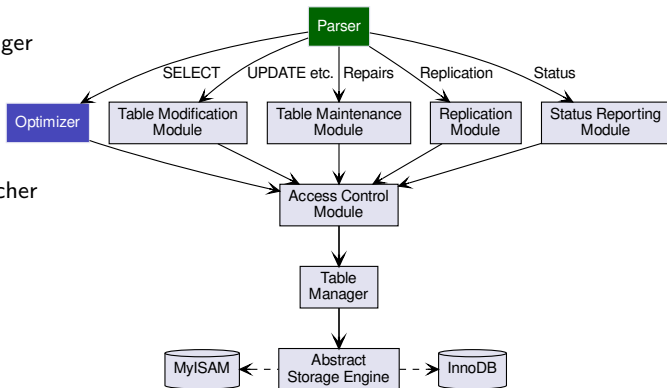
- ▶ ☒ 连接管理器 Connection Manager
- ▶ ☒ 线程管理器 Thread Manager
- ▶ ☒ 用户模块 User Module

## • 解析器

- ▶ ☒ 网络模块 Net Module
- ▶ ☒ 派发模块 Commander Dispatcher
- ▶ ☒ 词法分析 Lexical Analysis
- ▶ ☒ 语法分析 Syntax Analysis

## • 优化器

- ▶ ☒ 准备模块 Prepare Module
- ▶ ☒ 追踪日志 Optimizer Trace
- ▶ ☐ 优化模块 Optimize Module



2

## 子查询



# 子查询介绍

- 子查询 (Subquery) 是一种嵌套在另一个外部查询中的查询<sup>1</sup>

- 子查询有以下几种常见的形式

- ▶ 标量子查询 (Scalar Subquery) 指只返回一条数据的子查询

- ▶ 带有比较操作符号 (= <> > < >= <=) 的子查询

```
select * from t1 where c1 = (select max(c2) from t2);
```

- ▶ 带 IN/ANY/SOME/ALL 的子查询

```
select s1 from t1 where s1 = any(select s1 from t2);
```

```
select s1 from t1 where s1 in (select s1 from t2);
```

```
select s1 from t1 where s1 > all(select s1 from t2);
```

- ▶ 带 EXISTS/NOT EXISTS 的子查询

```
select * from t1 where exists (select * from t2);
```

```
select * from t1 where not exists (select * from t2);
```

- ▶ 派生表 (Derived Table) 是在查询过程中由子查询生成的虚拟表

```
select avg(sum1) from (select sum(c1) as sum1 from t1 group by c1) as t1;
```

- 子查询优化的示例<sup>2</sup>

- ▶ 例如：内部查询使用 union 合并，避免使用 or 连接两个子查询

<sup>1</sup><https://dev.mysql.com/doc/refman/8.0/en/subqueries.html>

<sup>2</sup><https://dev.mysql.com/doc/refman/8.0/en/optimizing-subqueries.html>



## 标量子查询 subq01.sql

```
mysql> select * from departments;
```

dept_no	dept_name
d009	Customer Service
d005	Development
d002	Finance
d003	Human Resources
d001	Marketing
d004	Production
d006	Quality Management
d008	Research
d007	Sales

9 rows in set (0.00 sec)

```
mysql> select * from dept_emp limit 3;
```

emp_no	dept_no	from_date	to_date
10001	d005	1986-06-26	9999-01-01
10002	d007	1996-08-03	9999-01-01
10003	d004	1995-12-03	9999-01-01

3 rows in set (0.01 sec)

- 采用子查询统计员工数，同时满足下面条件

- ▶ 员工号小于 10100
- ▶ 部门属于研发 (Development)

```
1  select
2      count(distinct emp_no) as emp_cnt
3  from
4      dept_emp a
5  where
6      a.emp_no < 10100
7      and a.dept_no = (
8          select
9              d.dept_no
10             from
11                 departments d
12             where
13                 d.dept_name = 'Development');
```





# 标量子查询分析

select#1 中的附属条件使用到标量表

```
{ "query_block": {
  "select_id": 1,
  "cost_info": {
    "query_cost": "6.69"
  },
  "table": {
    "table_name": "a",
    "access_type": "range",
    "possible_keys": ["PRIMARY", "dept_no"],
    "key": "dept_no",
    "used_key_parts": ["dept_no", "emp_no"],
    "key_length": "20",
    "rows_examined_per_scan": 32,
    "rows_produced_per_join": 32,
    "filtered": "100.00",
    "using_index": true,
    ...
    "used_columns": ["emp_no", "dept_no"],
    "attached_condition": "((`employees`.`a`.`dept_no` =
    /* select#2 */ select 'd005' from
      `employees`.`departments` `d` where true))
    and (`employees`.`a`.`emp_no` < 10100))",
```

access\_type 为 const 表示 select#2 是标量表 <sup>①</sup>

```
{ "attached_subqueries": [{
  "dependent": false,
  "cacheable": true,
  "query_block": {
    "select_id": 2,
    "cost_info": {
      "query_cost": "1.00"
    },
    "table": {
      "table_name": "d",
      "access_type": "const",
      "possible_keys": ["dept_name"],
      "key": "dept_name",
      "used_key_parts": ["dept_name"],
      "key_length": "162",
      "ref": ["const"],
      "rows_examined_per_scan": 1,
      "rows_produced_per_join": 1,
      "filtered": "100.00",
      "using_index": true,
      ...
```



## 唯一性的 IN 子查询 subq02.sql

```
mysql> select emp_no, first_name, last_name  
-> from employees limit 3;
```

emp_no	first_name	last_name
10001	Georgi	Facello
10002	Bezalel	Simmel
10003	Parto	Bamford

3 rows in set (0.00 sec)

```
mysql> select * from dept_manager limit 3;
```

emp_no	dept_no	from_date	to_date
110022	d001	1985-01-01	1991-10-01
110039	d001	1991-10-01	9999-01-01
110085	d002	1985-01-01	1989-12-17

3 rows in set (0.01 sec)

- 查询 d001 部门所有经理的姓名 <sup>❶</sup>

```
1 select  
2     e.first_name, e.last_name  
3 from  
4     employees e  
5 where  
6     e.emp_no in (  
7         select  
8             a.emp_no  
9         from  
10            dept_manager a  
11        where  
12            a.dept_no = 'd001');
```



# 非唯一性的 IN 子查询 subq03.sql

```
mysql> select * from departments limit 3;
```

dept_no	dept_name
d009	Customer Service
d005	Development
d002	Finance

3 rows in set (0.01 sec)

```
mysql> select * from dept_manager limit 3;
```

emp_no	dept_no	from_date	to_date
110022	d001	1985-01-01	1991-10-01
110039	d001	1991-10-01	9999-01-01
110085	d002	1985-01-01	1989-12-17

3 rows in set (0.01 sec)

- 查询 1995 年以来上任过经理的部门 <sup>1</sup>

```
1 select
2   *
3 from
4   departments d
5 where
6   d.dept_no in (
7     select
8       a.dept_no
9     from
10      dept_manager a
11    where
12      a.from_date >= '1995-01-01');
```



# IN 子查询转换半连接

- join\_preparation 阶段会将 IN (SELECT) 转换成 semi join
- subq02.sql 有唯一键走索引

```
1  select
2      `e`.`first_name` as `first_name`,
3      `e`.`last_name` as `last_name`
4  from
5      `employees` `e` semi join (`dept_manager` `a`)
6  where ((`a`.`dept_no` = 'd001') and (`e`.`emp_no` = `a`.`emp_no`))
```

- subq03.sql 非唯一键构造物化表

```
1  select
2      `d`.`dept_no` as `dept_no`,
3      `d`.`dept_name` as `dept_name`
4  from
5      `departments` `d` semi join (`dept_manager` `a`)
6  where ((`a`.`from_date` >= '1995-01-01')
7      and (`d`.`dept_no` = `a`.`dept_no`))
```



# EXISTS 子查询 subq04.sql

两个子查询的含义是一样的，最终都会变换成半连接进入 join\_optimization 阶段

- 带有 IN 的子查询

```
1  select
2      e.first_name,
3      e.last_name
4  from
5      employees e
6  where
7      e.emp_no in (
8          select
9              a.emp_no
10         from
11             dept_manager a
12         where
13             a.dept_no = 'd001');
14
```

- 带有 EXISTS 的子查询 ⓘ

```
1  select
2      e.first_name,
3      e.last_name
4  from
5      employees e
6  where
7      exists (
8          select
9              *
10         from
11             dept_manager a
12         where
13             a.dept_no = 'd001'
14             and a.emp_no = e.emp_no);
```



# EXISTS 子查询变换过程

- EXISTS 子查询通过 join\_preparation 转换成半连接，具体阶段如下：

- ▶ transformation\_to\_semi\_join
- ▶ transformations\_to\_nested\_joins

```
{  
  "transformation": {  
    "select#": 2,  
    "from": "IN (SELECT)",  
    "to": "semijoin",  
    "chosen": true,  
    "transformation_to_semi_join": {  
      "subquery_predicate": "exists(/* select#2 */ select 1 from `dept_manager` `a`  
        where ((`a`.`dept_no` = 'd001') and (`a`.`emp_no` = `e`.`emp_no`)))",  
      "embedded in": "WHERE",  
      "evaluating_constant_semijoin_conditions": [],  
      "semi-join condition": "((`a`.`dept_no` = 'd001') and (`e`.`emp_no` = `a`.`emp_no`))",  
      "decorrelated_predicates": [{"outer": "`e`.`emp_no`", "inner": "`a`.`emp_no`"}]},  
    }  
  }  
  "transformations_to_nested_joins": {  
    "transformations": [  
      "semijoin"  
    ],  
    "expanded_query": "/* select#1 */ select `e`.`first_name` AS `first_name`,  
      `e`.`last_name` AS `last_name` from  
      `employees` `e` semi join (`dept_manager` `a`)  
      where ((`a`.`dept_no` = 'd001') and (`e`.`emp_no` = `a`.`emp_no`))"}  
  }
```



3

## 半连接



# 半连接

- 半连接 (Semi-Join) 是一种数据库的查询方式
  - ▶ 半连接用来从两个或多个表中获取匹配的行
  - ▶ 但只返回第一个表中的行，而不包括第二个表中的任何列或重复的行
- MySQL 不支持直接使用半连接查询，但是以下两种查询会转换成半连接<sup>3</sup>
  - ▶ EXISTS 子查询
  - ▶ 包含 IN 的子句
- MySQL 提供了多种半连接优化策略<sup>4</sup>，这些策略通过优化器自动选择，具体如下
  - ▶ DuplicateWeedout (重复剔除)
  - ▶ FirstMatch (首次匹配)
  - ▶ LooseScan (松散扫描)
  - ▶ MaterializeLookup (索引式物化)
  - ▶ MaterializeScan (扫描式物化)

<sup>3</sup><https://dev.mysql.com/doc/refman/8.0/en/subquery-optimization.html>

<sup>4</sup><https://mariadb.com/kb/en/semi-join-subquery-optimizations/>





# 半连接举例

以表 t1 和 t2 进行半连接操作，其中 t1 是外表, t2 是内表

表 t1 和 t2 中的数据

```
mysql> select * from t1;
```

id	c
1	aaa
2	bbb
4	ccc
4	ddd

```
mysql> select * from t2;
```

id	c
1	xxx
3	yyy
4	uuu
4	vvv

- 内连接 innerjoin

```
mysql> select * from t1 inner join t2 using(id);
```

id	c	c
1	aaa	xxx
4	ddd	uuu
4	ccc	uuu
4	ddd	vvv
4	ccc	vvv

- 半连接 semijoin, 仅返回单表, 去重

```
mysql> select * from t1 semi join t2 using(id);
```

```
mysql> select * from t1 where t1.id in (select id from t2);
```

id	c
1	aaa
4	ddd
4	ccc

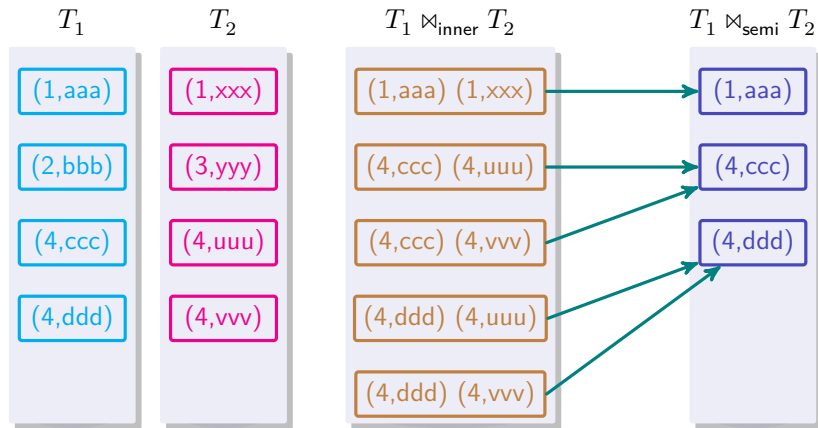
# 半连接算法

MySQL 提供了多种半连接优化策略，这些策略通过优化器自动选择，具体如下

- ① DuplicateWeedout (重复剔除)
  - ▶ 执行普通的两表内连接操作，用临时表缓存结果
  - ▶ 并在临时表中通过主键或唯一索引去除重复的元组
- ② FirstMatch (首次匹配)
  - ▶ 扫描内部表的行组合
  - ▶ 当有多个实例的给定值组时，只选择一个而不是全部返回
  - ▶ 这种快速扫描并消除了不必要的行的生成
- ③ LooseScan (松散扫描)
  - ▶ 在执行连接时，如果内表的元组有序（通常通过索引实现）
  - ▶ 则可以根据索引拿出每组重复元组中的第一个元组与外表进行连接，跳过后续相同的记录
- ④ MaterializeLookup (索引式物化)
  - ▶ 将子查询的结果物化到临时表，并为其创建索引
  - ▶ 执行连接时，使用临时表的索引来完成连接操作
- ⑤ MaterializeScan (扫描式物化)
  - ▶ 类似于 MaterializeLookup，但临时表的索引不能辅助加快连接
  - ▶ 只能通过全表扫描的方式完成半连接操作

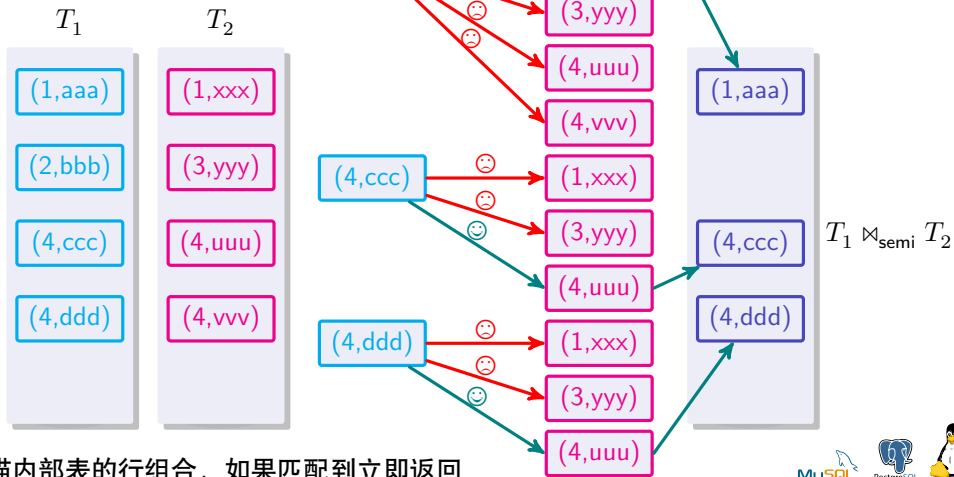


# DuplicateWeedout

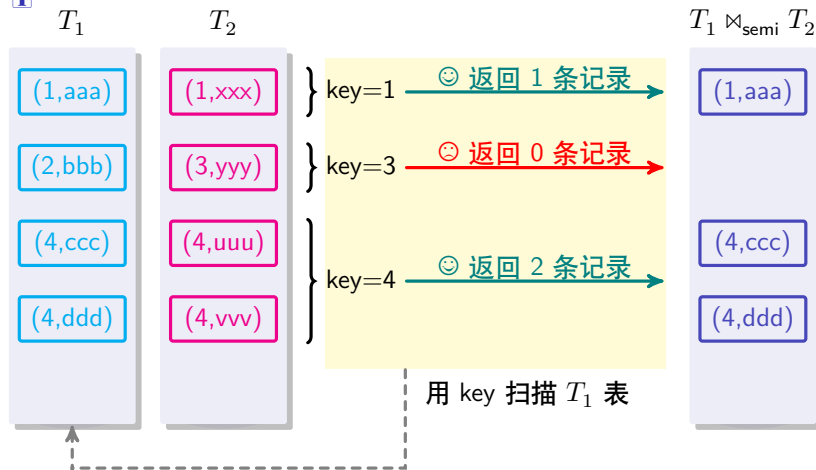


- 1 先计算表  $T_1$  和  $T_2$  的内连接  $T_1 \bowtie_{\text{inner}} T_2$ ，放入临时表
- 2 将临时表去重得到最终结果半连接表  $T_1 \bowtie_{\text{semi}} T_2$

# FirstMatch

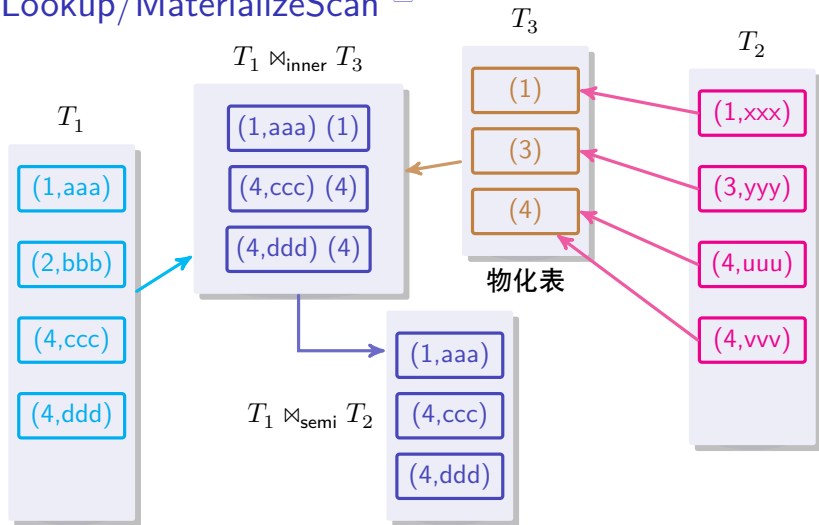


- 扫描内部表的行组合，如果匹配到立即返回



- ① 在执行连接时，如果内表的元组有序（通常通过索引实现）
- ② 则可以根据索引拿出每组重复元组中的第一个元组与外表进行连接
- ③ 跳过后续相同的记录

# MaterializeLookup/MaterializeScan



- 将  $T_2$  物化到临时表  $T_3$ ，然后计算  $T_1$  内连接  $T_3$  得到结果
- MaterializeLookup 对  $T_2$  创建索引，MaterializeScan 不建索引直接扫描

# 不同策略 Cost 值对比

- 通过 hints 可以控制选取对于的半连接策略<sup>5</sup>
- 下面是两个查询优化值的对比

- subq10 查询 ⓘ

```
select * from departments d where  
  d.dept_no in (  
    select a.dept_no  
    from dept_manager a  
    where a.from_date >= '1995-01-01');
```

- subq10 策略对比

策略	Cost	说明
DuplicateWeedout	8.05	
FirstMatch	8.95	
LooseScan	-	不支持
Materialize	6.50	命中 ☺

- subq20 查询 ⓘ

```
select * from employees e where  
  e.emp_no in (  
    select a.emp_no  
    from dept_manager a  
    where a.dept_no >= 'd003')
```

- subq20 策略对比

策略	Cost	说明
DuplicateWeedout	16.59	
FirstMatch	124440.10	代价太大
LooseScan	10.39	命中 ☺
Materialize	16.59	

<sup>5</sup><https://dev.mysql.com/doc/refman/8.0/en/optimizer-hints.html>



# 优化器配置开关

- 优化器配置开关默认值通过数据库系统变量来控制

```
mysql> show variables like '%optimizer_switch%'\G
```

```
***** 1. row *****
```

```
Variable_name: optimizer_switch
```

```
Value: index_merge=on, index_merge_union=on, index_merge_sort_union=on,  
       index_merge_intersection=on, engine_condition_pushdown=on,  
       index_condition_pushdown=on, mrr=on, mrr_cost_based=on,  
       block_nested_loop=on, batched_key_access=off, materialization=on,  
       semijoin=on, loosescan=on, firstmatch=on, duplicateweedout=on,  
       subquery_materialization_cost_based=on, use_index_extensions=on,  
       condition_fanout_filter=on, derived_merge=on, use_invisible_indexes=off,  
       skip_scan=on, hash_join=on, subquery_to_derived=off,  
       prefer_ordering_index=on, hypergraph_optimizer=off,  
       derived_condition_pushdown=on
```

```
1 row in set (0.02 sec)
```

- 可以使用下面方式开启/关闭

```
set optimizer_switch='semijoin=on'
```

```
set optimizer_switch='semijoin=off'
```





# 半连接优化函数

- 半连接优化决策函数 `advance_sj_state()`，计算不同半连接策略的 `cost` 值

► `sql/sql_planner.cc`

```
4105 void Optimize_table_order::advance_sj_state(table_map remaining_tables,
      :
4597     if (sj_strategy != SJ_OPT_NONE)
4598         pos->set_prefix_cost(best_cost, best_rowcount);
4599 }
```

- 半连接优化计算伪代码, 对每种策略分别计算, 最后保存最优

```
for each semi-join strategy {
    update strategy's state variables;
    if (join prefix has all the tables that are needed to consider
        using this strategy for the semi-join(s)) {
        calculate cost of using the strategy
        if ((this is the first strategy to handle the semi-join nest(s) ||
            the cost is less than other strategies)) {
            // Pick this strategy
            pos->sj_strategy= ..
            ..
        }
    }
}
```



4

## 其他类型



# NOT IN 子查询 subq05.sql

```
mysql> select * from departments limit 3;
```

dept_no	dept_name
d009	Customer Service
d005	Development
d002	Finance

3 rows in set (0.01 sec)

```
mysql> select * from dept_manager limit 3;
```

emp_no	dept_no	from_date	to_date
110022	d001	1985-01-01	1991-10-01
110039	d001	1991-10-01	9999-01-01
110085	d002	1985-01-01	1989-12-17

3 rows in set (0.01 sec)

- 使用 NOT IN 的查询语句会被变化成 antijoin <sup>1</sup>

```
1 select
2   *
3 from
4   departments d
5 where
6   d.dept_no not in (
7     select
8       a.dept_no
9     from
10      dept_manager a
11    where
12      a.from_date > '1995-01-01');
```



## 派生表 subq06.sql

```
mysql> select emp_no, first_name, last_name  
-> from employees limit 3;
```

emp_no	first_name	last_name
10001	Georgi	Facello
10002	Bezalel	Simmel
10003	Parto	Bamford

3 rows in set (0.00 sec)

```
mysql> select * from titles limit 3;
```

emp_no	title	from_date	to_date
10001	Senior Engineer	1986-06-26	9999-01-01
10002	Staff	1996-08-03	9999-01-01
10003	Senior Engineer	1995-12-03	9999-01-01

3 rows in set (0.01 sec)

- 使用派生表使用 merge 策略来优化 ⓘ

```
1  select  
2      *  
3  from  
4      titles t,  
5      (  
6          select  
7              *  
8          from  
9              employees  
10             where  
11                 emp_no <= 10100) emp_100  
12  where  
13      t.emp_no = emp_100.emp_no  
14  and t.from_date > '2000-01-01';
```



# 结束

