

第十九讲：性能分析及 handleron 存储引擎接口设计

知春路遇上八里桥

<2024-07-22 Mon>



- 1 前情提要
- 2 追溯日志
- 3 性能分析
- 4 PFS 实现
- 5 存储引擎的接口层




1

前情提要



本节内容



- 连接器

- ▶ ☒ 连接管理器 Connection Manager
- ▶ ☒ 线程管理器 Thread Manager
- ▶ ☒ 用户模块 User Module

- 解析器

- ▶ ☒ 网络模块 Net Module
- ▶ ☒ 派发模块 Commander Dispatcher
- ▶ ☒ 词法分析 Lexical Analysis
- ▶ ☒ 语法分析 Syntax Analysis

- 优化器

- ▶ ☒ 准备模块 Prepare Module
- ▶ ☒ 追踪日志 Optimizer Trace
- ▶ ☒ 优化模块 Optimize Module

- 执行器

- ▶ ☒ 火山模型 Volcano Model
- ▶ ☒ 执行模块 Execution Module

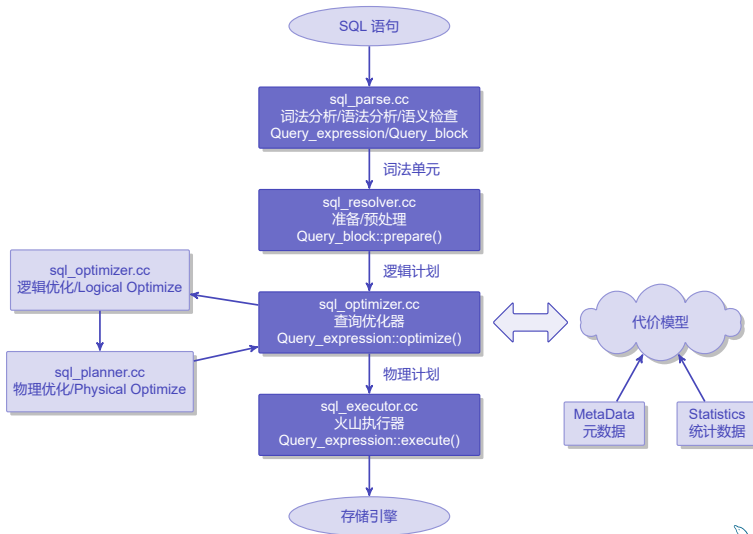


2

追溯日志



服务层架构梳理



Parser 解析器

- `parse_sql()` 函数进入 SQL 解析器
 - ▶ 词法分析, 产生 token
 - ▶ 语法分析, 产生抽象语法树
- 通过 BISON 生成解析树 PT
 - ▶ Parse Tree
- 语法分析还构建词法单元 LU
 - ▶ Lexical Unit
 - ▶ `Query_expression` 表示查询表达式
 - ▶ `Query_block` 表示查询块

```
T@10: | | | >parse_sql
T@10: | | | | >THD::convert_string
...
T@10: | | | | >MEM_ROOT::AllocSlow
T@10: | | | | | enter: root: 0x7fff300aed10
T@10: | | | | | >MEM_ROOT::AllocBlock
T@10: | | | | | | >char* PFS_instr_name::str
T@10: | | | | | | <char* PFS_instr_name::str
T@10: | | | | | <MEM_ROOT::AllocBlock
T@10: | | | | | <MEM_ROOT::AllocSlow
T@10: | | | | | >THD::convert_string
T@10: | | | | | >find_udf
...
T@10: | | | | | >Query_block::add_table_to_list
T@10: | | | | | >MEM_ROOT::AllocSlow
T@10: | | | | | | enter: root: 0x7fff300aed10
T@10: | | | | | | >MEM_ROOT::AllocBlock
T@10: | | | | | | | >char* PFS_instr_name::str
T@10: | | | | | >Query_block::add_joined_table
T@10: | | | | <parse_sql
```



Resolver 准备模块

```
T010: | | | >mysql_execute_command
...
T010: | | | | >bool Sql_cmd_dml::execute
T010: | | | | >bool Sql_cmd_dml::prepare
T010: | | | | | >check_table_access
T010: | | | | | >open_tables_for_query
...
T010: | | | | | >Query_block::prepare
T010: | | | | | >Query_block::setup_tables
T010: | | | | | >setup_fields
T010: | | | | | >Query_block::setup_conds
T010: | | | | | >Query_block::apply_local_transforms
T010: | | | | | | >Query_block::simplify_joins
T010: | | | | | | <Query_block::simplify_joins
T010: | | | | | | >build_bitmap_for_nested_joins
T010: | | | | | | <build_bitmap_for_nested_joins
T010: | | | | | | <Query_block::apply_local_transforms
T010: | | | | | | opt: steps: ending struct
T010: | | | | | | opt: join_preparation: ending struct
T010: | | | | | | opt: (null): ending struct
T010: | | | | | | <Query_block::prepare
T010: | | | | | <bool Sql_cmd_dml::prepare
```

- 准备阶段是预处理词法单元
 - ▶ mysql_execute_command() 执行命令
 - ▶ Sql_cmd_dml::execute() SELECT 语句
- 各种 prepare 函数叠加调用
 - ▶ Sql_cmd_dml::prepare()
 - ▶ Query_block::prepare()
- 单个查询块的处理包括
 - ▶ setup_tables() 设置表
 - ▶ setup_fields() 设置列
 - ▶ setup_conds() 设置条件
 - ▶ apply_local_transforms() 本地变换
- 准备阶段生成结果是逻辑计划
 - ▶ Logical Plan
- 准备阶段对应 Opt_trace 的阶段是
 - ▶ join_preparation
 - ▶ 注意 join_preparation 可能会嵌套



Optimizer 优化器

- `Query_expression::optimize()` 是优化器的入口, 最终调用 `JOIN::optimize()` 执行优化
- 优化器核心目的是计算不同策略的 `cost` 值, 最终选取 `cost` 值最小的执行计划
- 在计算 `JOIN` 次序时如果搜索使用次序排列效率较低, 故使用贪心算法 `greedy_search()`

```
T@10: | | | | | >Query_expression::optimize
T@10: | | | | | >Query_block::optimize
T@10: | | | | | >JOIN::optimize
...
T@10: | | | | | >optimize_cond
T@10: | | | | | >JOIN::make_join_plan
T@10: | | | | | >int ha_innabase::info_low
T@10: | | | | | >JOIN::make_sum_func_list
T@10: | | | | | >get_quick_record_count
T@10: | | | | | >test_quick_select
...
T@10: | | | | | >Optimize_table_order::greedy_search
T@10: | | | | | >Optimize_table_order::best_extension_by_limited_search
part_plan; idx: 0 best: DBL_MAX atime: 0 itime: 0 count: 1
POSITIONS:
BEST_REF: employees(9,276023,1)
T@10: | | | | | >Optimize_table_order::best_access_path
...
T@10: | | | | | opt: cost: 2.81809
T@10: | | | | | opt: chosen: 1
...
```



Executor 执行器

- `Query_expression::execute()` 是执行器的入口, 处理执行计划并从存储引擎中读取数据
- 执行器采用火山模型, 执行器对应 `Opt_trace` 的 `join_execution` 阶段

```
T@10: | | | | | >Query_expression::execute
T@10: | | | | | >THD::send_result_metadata
T@10: | | | | | >bool Protocol_classic::start_result_metadata
T@10: | | | | | >bool Protocol_classic::send_field_metadata
T@10: | | | | | >bool Protocol_classic::end_row
...
T@10: | | | | | >InitIndexRangeScan
T@10: | | | | | >handler::ha_index_init
T@10: | | | | | >int ha_innbase::index_init
T@10: | | | | | >ha_innbase::change_active_index
T@10: | | | | | >dict_index_t* ha_innbase::innbase_get_index
... 读取/发送数据若干次
T@10: | | | | | >int IndexRangeScanIterator::Read
T@10: | | | | | >handler::ha_multi_range_read_next
T@10: | | | | | >int handler::multi_range_read_next
T@10: | | | | | >int handler::read_range_first
T@10: | | | | | >handler::ha_index_first
T@10: | | | | | >int ha_innbase::index_first
T@10: | | | | | >int ha_innbase::index_read
T@10: | | | | | >row_search_mvcc
...
T@10: | | | | | >bool Query_result_send::send_data
T@10: | | | | | >THD::send_result_set_row
...
```



3

性能分析



Profile 配置

```
mysql> set profiling=1;
Query OK, 0 rows affected, 1 warning (0.00 sec)
```

```
mysql> \. exe02.sql
... 省略执行结果输出
```

```
mysql> show profiles\G
***** 1. row *****
Query_ID: 1
Duration: 0.00713825
  Query: select
    x.dept_no,
    e.first_name,
    e.last_name
  from
    dept_manager x
  left join
    employees e
    using (emp_no)
  where
    x.dept_no < 'd004'
1 row in set, 1 warning (0.01 sec)
```

- 开启 profiling 后执行查询语句
- 通过 show profiles 查看 Query_ID
show profile cpu, ipc for query 1;
show profile memory for query 1;

```
mysql> show profile for query 1;
+-----+-----+
| Status                               | Duration |
+-----+-----+
| starting                             | 0.000921 |
| checking permissions                  | 0.000066 |
| checking permissions                  | 0.000135 |
| Opening tables                        | 0.007945 |
| end                                   | 0.000027 |
| query end                             | 0.000042 |
| closing tables                        | 0.000036 |
| freeing items                         | 0.000117 |
| cleaning up                           | 0.000057 |
+-----+-----+
9 rows in set, 1 warning (0.00 sec)
```



Profile 结果查询和分析

- 通过 Query_ID 查询 information_schema.profilng 表获取执行耗时

```
mysql> select query_id, seq, state, duration, source_file, source_line  
-> from information_schema.profilng where query_id = 1;
```

query_id	seq	state	duration	source_file	source_line
1	2	starting	0.000692	NULL	NULL
1	3	Executing hook on transaction	0.000054	rpl_handler.cc	1477
1	4	starting	0.000142	rpl_handler.cc	1479
1	5	checking permissions	0.000041	sql_authorization.cc	2146
1	6	checking permissions	0.000039	sql_authorization.cc	2146
1	7	Opening tables	0.000697	sql_base.cc	5797
1	8	init	0.000061	sql_select.cc	772
1	9	System lock	0.000171	lock.cc	332
1	10	optimizing	0.000171	sql_optimizer.cc	354
1	11	statistics	0.001304	sql_optimizer.cc	694
1	12	preparing	0.000531	sql_optimizer.cc	778
1	13	executing	0.002607	sql_union.cc	1670
1	14	end	0.000042	sql_select.cc	805
1	15	query end	0.000024	sql_parse.cc	4881
1	16	waiting for handler commit	0.000151	handler.cc	1610
1	17	closing tables	0.000190	sql_parse.cc	4944
1	18	freeing items	0.000164	sql_parse.cc	5413
1	19	cleaning up	0.000060	sql_parse.cc	2489

18 rows in set, 1 warning (0.01 sec)



performance_schema

- `show profiles` 在后续 MySQL 中逐渐弃用，取而代之的是 `performance_schema`

'SHOW PROFILES' is deprecated and will be removed in a future release.
Please use Performance Schema instead

- 在配置文件中启用 `performance_schema` 参数，打开后将会收集所有用户的执行历史事件

```
[mysqld]  
performance_schema=ON
```

- `performance_schema` 提供了多种表来存储监控数据，具体分为以下几类

- ▶ Setup 表：用于配置 `performance_schema` 的行为，比如哪些监控器被启用
 - ① `setup_instruments` 表来启用或禁用特定的监控器 (instruments)
- ▶ Instances 表：包含有关监控实例（如线程、表等）的信息，例如：xxx_instances
- ▶ Events 表：包含关于 MySQL 服务器事件（如语句执行、等待事件等）的信息，例如：
 - ① `events_waits_summary_by_instance`：提供关于等待事件的汇总信息
 - ② `events_statements_summary_by_digest`：提供关于 SQL 语句执行情况的汇总信息，按语句摘要分组
 - ③ `events_transactions_summary_by_state`：提供关于事务状态的汇总信息
- ▶ Consumer 表：定义了如何消费（存储或汇总）`performance_schema` 中的数据



performance_schema 参数配置

❶ 配置 setup_actors, 开启监控的用户

```
select * from performance_schema.setup_actors;
```

❷ 配置 setup_instruments, 开启对 'statements' 和 'stage' 名称的配置

```
update performance_schema.setup_instruments set ENABLED = 'YES', TIMED = 'YES'  
  where name like '%statement/%';  
update performance_schema.setup_instruments set ENABLED = 'YES', TIMED = 'YES'  
  where name like '%stage/%';
```

❸ 配置 setup_consumer, 启用 events_statements_*, events_stages_* 开头的事件类型消费

```
update performance_schema.setup_consumers set ENABLED = 'YES'  
  where name like '%events_statements_%';  
update performance_schema.setup_consumers set ENABLED = 'YES'  
  where name like '%events_stages_%';
```



performance_schema 执行并收集结果

- 查询事件的 EVENT_ID

```
select EVENT_ID,  
       truncate(TIMER_WAIT / 1000000000000, 6) as Duration,  
       SQL_TEXT  
from performance_schema.events_statements_history_long  
where SQL_TEXT like '%empl%'\G
```

- 查询耗时结果

```
select  
  event_name as Stage,  
  truncate(TIMER_WAIT / 1000000000000, 6) as Duration  
from performance_schema.events_stages_history_long  
where NESTING_EVENT_ID = 422;
```



performance_schema 查询收集结果

```
mysql> select event_name as Stage,  
-> truncate(TIMER_WAIT / 1000000000000, 6) as Duration  
-> from performance_schema.events_stages_history_long  
-> where NESTING_EVENT_ID = 422;
```

Stage	Duration
stage/sql/starting	0.0010
stage/sql/Executing hook on transaction begin.	0.0000
stage/sql/starting	0.0001
stage/sql/checking permissions	0.0000
stage/sql/checking permissions	0.0000
stage/sql/Opening tables	0.0011
stage/sql/init	0.0000
stage/sql/System lock	0.0002
stage/sql/optimizing	0.0002
stage/sql/statistics	0.0018
stage/sql/preparing	0.0006
stage/sql/executing	0.0028
stage/sql/end	0.0000
stage/sql/query end	0.0000
stage/sql/waiting for handler commit	0.0001
stage/sql/closing tables	0.0002
stage/sql/freeing items	0.0001
stage/sql/cleaning up	0.0000

18 rows in set (0.01 sec)

```
mysql> select EVENT_ID,  
-> truncate(TIMER_WAIT / 1000000000000, 6)  
-> as Duration,  
-> SQL_TEXT  
-> from  
-> events_statements_history_long  
-> where SQL_TEXT like '%empl%';  
...  
***** 2. row *****  
EVENT_ID: 422  
Duration: 0.0091  
SQL_TEXT: select  
          x.dept_no,  
          e.first_name,  
          e.last_name  
from  
  dept_manager x  
  left join  
    employees e  
    using (emp_no)  
where  
  x.dept_no < 'd004'
```

4

PFS 实现



performance_schema

- performance_schema¹ 的实现主要位于 `storage/perfschema/` 目录下
 - ▶ 该目录下有很多代码, 这些函数一般是 `PSI_xxx` 或 `PFS_xxx` 作为前缀
- performance_schema 使用了多种数据结构来存储和管理监控数据
 - ▶ 包括哈希表、链表、位图等
 - ▶ 这些数据结构的选择和实现在很大程度上决定了 performance_schema 的性能和可扩展性
- 事件的采集通常是通过在 MySQL 代码的关键位置插入钩子 (hooks) 来实现的
 - ▶ 这些钩子在特定事件发生时被触发
 - ▶ 调用 performance_schema 的相关函数来记录事件信息
 - ▶ 例如在之前说的在线程创建是有 PFS 的钩子 `storage/perfschema/pfs.cc`


```
3002 extern "C" {  
3003     static void *pfs_spawn_thread(void *arg) {  
3004         auto *typed_arg = (PFS_spawn_thread_arg *)arg;  
3005         void *user_arg;  
3006         void *(*user_start_routine)(void *);
```

- 收集到的数据后, 用户可以根据需求进行汇总和整理
 - ▶ performance_schema 提供了 SQL 查询接口, 以便用户可以方便地查询和分析

¹https://dev.mysql.com/doc/dev/mysql-server/8.0.37/PAGE_PFS.html



setup_instruments 表创建

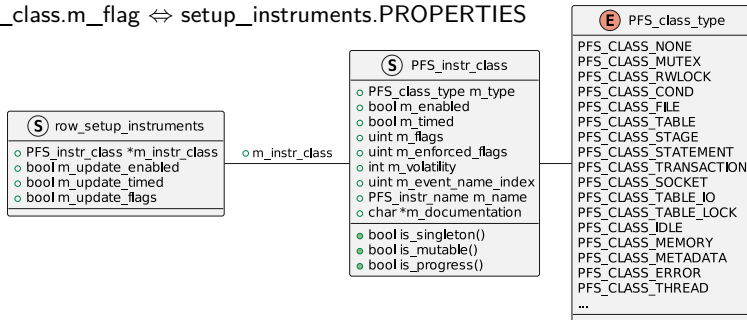
- 表定义代码  storage/perfschema/table_setup_instruments.cc

```
19 Plugin_table table_setup_instruments::m_table_def(  
20     /* Schema name */  
21     "performance_schema",  
22     /* Name */  
23     "setup_instruments",  
24     /* Definition */  
25     " NAME VARCHAR(128) not null,\n"  
26     " ENABLED ENUM ('YES', 'NO') not null,\n"  
27     " TIMED ENUM ('YES', 'NO'),\n"  
28     " PROPERTIES SET('singleton', 'progress', 'user', 'global_statistics', "  
29     "'mutable', 'controlled_by_default') not null,\n"  
30     " FLAGS SET('controlled'),\n"  
31     " VOLATILITY int not null,\n"  
32     " DOCUMENTATION LONGTEXT,\n"  
33     " PRIMARY KEY (NAME) USING HASH\n",  
34     /* Options */  
35     " ENGINE=PERFORMANCE_SCHEMA",  
36     /* Tablespace */  
37     nullptr);
```



setup_instruments 表数据结构

- setup_instruments² 表中的行数据主要是通过 row_setup_instruments 类来存取的
- 元数据主要存储在 pfs_instr_class 对象中
- pfs_instr_class 对象描述了 instrument 的信息, 该类和 setup_instruments 表对应如下:
 - ▶ pfs_instr_class.m_name ⇔ setup_instruments.NAME
 - ▶ pfs_instr_class.m_enabled ⇔ setup_instruments.ENABLED
 - ▶ pfs_instr_class.m_flag ⇔ setup_instruments.PROPERTIES
 - ▶ ...



²<https://dev.mysql.com/doc/refman/8.0/en/performance-schema-setup-instruments-table.html>

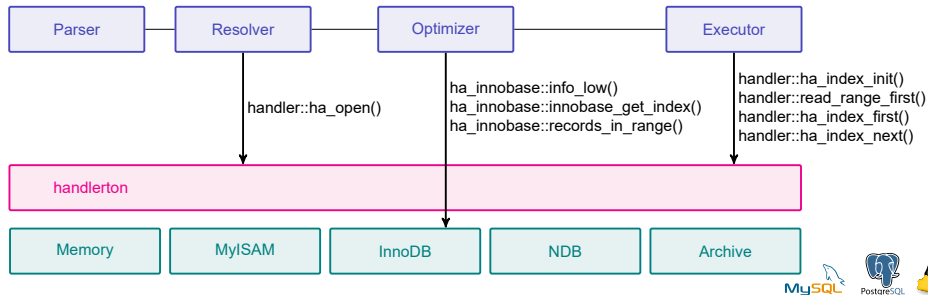
5

存储引擎的接口层



handlerton

- handlerton 代表了 MySQL 存储引擎的接口层, (handlerton = table handler + singleton)
 - ▶ 存储引擎早期就叫表处理器 (table handler)
 - ▶ handlerton 是一个单例 (singleton) 的数据结构, 每个存储引擎只用一个实例
 - ▶ handlerton 结构就是 MySQL 服务器与这些存储引擎之间交互的桥梁
 - ▶ 可以通过实现 handlerton 的接口来创建自己的存储引擎, 并通过插件方式嵌入到 MySQL 中
- MySQL 支持多种存储引擎, 每种存储引擎都提供了不同特性
 - ▶ 例如: 存储机制、锁定级别、事务支持等
 - ▶ 常见的存储引擎包括: InnoDB, MyISAM



handlerton 实现

- handlerton³ 包含了以下几个关键部分
 - ▶ 状态信息：如存储引擎的名称、是否支持事务、是否支持外键等
 - ▶ 函数指针：指向存储引擎实现的函数
 - ① 从表的创建、删除、打开、关闭
 - ② 行的插入、删除、更新、读取
 - ③ 索引的创建、删除、查询等
 - ▶ 代码中实现 handler 抽象类, 使用 ha_ 作为大部分函数名称前缀
- handlerton 在 `sql/handler.h` 中声明, 在 `sql/handler.cc` 中实现

```
2621 struct handlerton {
2622     /**
2623      * Historical marker for if the engine is available or not.
2624      */
2625     SHOW_COMP_OPTION state;
2626
2627     /**
2628      * Historical number used for frm file to determine the correct storage engine.
2629      * This is going away and new engines will just use "name" for this.
2630      */
2631     enum legacy_db_type db_type;
```

³<https://dev.mysql.com/doc/dev/mysql-server/8.0.37/structhandlerton.html>



handlernton 的实例 xxx_hton

- handlernton 的实例通常静态定义在 ha_xxx.cc 中, 服务层通过函数指针来调用

// 通常定义成

```
static handlernton { ... } xxx_hton;
```

// 具体的 hton 实例例如

```
extern handlernton *myisam_hton;
```

```
extern handlernton *heap_hton;
```

```
extern handlernton *temptable_hton;
```

```
extern handlernton *innodb_hton;
```

- 例如: InnoDB 的 hton 实例, 可以通过 gdb 打印其内容

```
(gdb) p *innodb_hton
```

```
$10 = {
```

```
    state = SHOW_OPTION_YES,
```

```
    db_type = DB_TYPE_INNODB,
```

```
    slot = 1,
```

```
    savepoint_offset = 0,
```

```
    close_connection = 0x55555a661380 <innobase_close_connection(handlernton*, THD*)>,
```

```
    kill_connection = 0x55555a66181f <innobase_kill_connection(handlernton*, THD*)>,
```

```
    pre_dd_shutdown = 0x55555a654ba1 <innodb_pre_dd_shutdown(handlernton*)>,
```

```
    savepoint_set = 0x55555a661170 <innobase_savepoint(handlernton*, THD*, void*)>,
```

```
    savepoint_rollback = 0x55555a660c46 <innobase_rollback_to_savepoint(handlernton*, THD*, void*)>,
```

```
    ...
```



结束

